

Try one bounded H.A.R.D. review

H.A.R.D. Protocol 0.3 · Public Preview

Human-centered AI Readiness and Decision Protocol

The invitation

Try the guide with one plan, prototype or implementation you know. We want to learn whether it helps you identify the decisions inside the artifact, find the relevant evidence and understand what to do next. Report unclear questions and incomplete reviews as well as useful findings.

Participation is voluntary. This is formative use, not certification, an endorsement request or a controlled effectiveness study. Do not send confidential client, customer, participant or health information. Agree retention and permitted disclosure before sharing.

Pick the route that fits your task

For a solo practitioner or small team, begin with `artifact_review`. A plan without code is in scope. State its revision, intended use, current scope and review stage. No measured baseline, ROI or independently staffed evaluation is required.

Inspect one consequential choice: what was selected, what purpose it serves, the alternatives, reasons, assumptions, tradeoffs, evidence and human decision still needed. Explicitly record whether engineering deepening applies and why. If it does, inspect only the relevant truth, ownership, state, boundary/contract, failure, time/ordering and assumption surfaces and name one bounded condition that could disconfirm the choice. Separate documented past reasons from current new proposals. Keep the original choice when justified. Record specified, walkthrough, implemented and runtime-tested evidence separately for the current requirements and recovery paths.

If your goal is funding engineering, use QUICK6/FULL with all six gates. QUICK6 is for low risk with evidence available; its ≤ 15 -minute target is untested. Higher or unknown risk requires FULL. Preserve missing evidence and stopped gates. A minimum artifact review does not mean these gates passed.

If your goal is measuring reviewer performance, first plan `independent_evaluation` with an eligible frozen reference, roles and locked judgments. A solo or agent-assisted review can have useful findings even while those metrics are N/A.

Share only what you choose to permit

- Exact versions, artifact kind/revision, stage, population, route, role overlap and whether criteria were set before review.
- Preparation and session effort if actually recorded, confusing questions and evidence that was hard to find.
- Choices kept or questioned, engineering-deepening disposition, relevant decision surfaces/challenges, evidence gaps, stage result or engineering gates, next action and owner.
- Ineligible metrics and the reason for each N/A. Do not invent before/after judgments or timings retrospectively.
- Whether you understood what the result allows and what remains undecided.
- Your disclosure choice: private, permission-checked anonymous summary or separately approved attribution.

A facilitator may retain source locations or fingerprints without publishing the source itself. You may stop at any time. An incomplete review remains useful feedback; do not fill unknowns to obtain a favorable result.

What follows

With permission, the observations may inform revisions. The maintainer distinguishes reported experience from independently inspected evidence. One use cannot demonstrate effectiveness, safety, defect prevention or general adoption. Public material will not include your identity or quotes without the appropriate permission.

Start: <https://www.takyejun.com/research/ai-readiness>
Updates: <https://www.takyejun.com/research/ai-readiness/updates>